

A decorative graphic on the left side of the slide consists of two overlapping parallelograms. The front one is blue and the back one is a light green. They are positioned diagonally, with the blue one partially covering the green one.

CSEN 140L - Lab 8

Nolan Anderson



Lab 8 High-Level Objectives

- Conduct K-Means Clustering on the MNIST dataset
- Complete the tasks on the lab assignment PDF a single Jupyter notebook (.ipynb) file
- Use only the **NumPy** and **Matplotlib** libraries for your implementation
 - No machine learning libraries (except for the data retrieval part given by professor)
- Demo to me and submit your notebook to Camino before leaving the lab section



The Data Explained

- The MNIST dataset consists of 28 x 28 grayscale images of handwritten digits (0 – 9).
 - Since the images are 28 x 28, each image has 784 pixels.
 - Each pixel counts as a feature.
 - Because the dataset is quite large, we will just be using the first 2000 samples.
 - As such, your subset of the data should be (2000 x 784).



Clustering and K-Means

- Clustering is unsupervised machine learning algorithm.
 - Because of this, we don't have access to the true labels during the training.
- K-Means is an example of a clustering algorithm.
 - K-Means has a hyperparameter K , the number of clusters.
 - In this case, since we know there are 10 types of digits, we require you to set $K = 10$, but strictly speaking the value of K is your choice when clustering.



K-Means Algorithm

- Initialize K centroids
 - In our case, we are going to randomly select K samples from our dataset to be our initial centroids.
- Optimize the clusters
 - Calculate the distance of all samples to the K centroids
 - Assign samples to the closest centroids
 - Calculate a new centroid for each of the K clusters
 - Repeat the above optimization steps until convergence
 - Convergence can be defined in terms of cluster assignments (i.e. no changes in cluster assignments) or centroids (i.e. no changes in the centroids)



Euclidean Distance

- When we refer to distance in the K-Means Algorithm, we are talking about euclidean distance:

$$\sum_1^N [(X_i - c)^2]$$

- X_i is the i th sample in the dataset (N total)
- c is a centroid



K-Means Loss Function

- For clustering, we can use the sum of squared error (SSE) as our loss function.
- For each cluster, SSE is computed as follows:

$$SSE = \sum_1^N (X_i - \bar{X})^2$$

- X_i = i th sample of the cluster
 - \bar{X} = the cluster's centroid
-
- The total SSE for your clustering results is simply the sum of the SSE values from each cluster.



General Code Structure

- `def my_kmeans(X, K, M):`
 - `// outermost loop cycles over different initializations`
 - `for m in range(M):`
 - `centroids = randomly selected each time`
 - `// actual kmeans loop below`
 - `while (not converged):`
 - `for k in range(K):`
 - `compute distance between samples X and centroid[k]`
 - `assign observations to clusters (use argmin)`
 - `for k in range(K):`
 - `recompute centroid[k] by taking the mean of each cluster`
 - `Check convergence and stop loop if converged`



Tips

- Use a fixed random seed in your implementation to allow for repeatability
 - See the next slide for a suggestion of how to do this
- Use the *plot_centroids* function provided in the Files to plot your final centroids to see if they are reasonable
 - I have provided two sample centroids in that file so you can see what the plot is supposed to look like
- Don't forget that you are required to use $K = 10$ and $M = 15$



Helpful NumPy Functions

- `rng = np.random.default_rng(SEED)`
 - Allows you to instantiate a random number generator with a fixed seed (*SEED*)
 - `rng` has a method *choice* that allows you to randomly select a subset of a given set (useful for initial centroid selection)
- `np.argmin()`
 - Useful to find cluster assignments
- `np.where()`
 - Useful to select samples belonging to a particular cluster
- `np.mean()`
 - Useful for recomputing the centroids